

Project 1: Parallel Addition of Sparse Matrices using PThreads

Let A and B be two *sparse* $m \times n$ matrices, i.e., matrices most elements of which are zero; let A and B be stored in the so-called *Yale Sparse Matrix Format (YSMF)*. Besides sparsity, there is no specific property or structure to be assumed for the matrices.

Write and test a parallel program using PThreads that computes $C = A + B$, where C is again a sparse $m \times n$ matrix and is to be stored in the YSMF. The number of threads p that your parallel program deploys, should be easily configurable.

Hints and explanations:

1. On sparse matrices and YSMF (adapted from: *Patterson, Hennessy. Computer Organization and Design. Fourth Edition. Morgan Kaufmann 2009*; but almost identically also to be found in Wikipedia):

Storing a sparse matrix

Many if not most entries of a sparse matrix are zeros. The basic idea when storing sparse matrices is to store only the non-zero entries as opposed to storing all entries. Depending on the number and distribution of the non-zero entries, different data structures can be used and yield huge savings in memory when compared to a naive approach.

One example of such a sparse matrix format is the Yale Sparse Matrix Format. It stores an initial sparse $m \times n$ matrix, M , in row form using three one-dimensional arrays. Let NNZ denote the number of nonzero entries of M . The first array is A , which is of length NNZ , and holds all nonzero entries of M in left-to-right top-to-bottom order. The second array is IA , which is of length $m + 1$ (i.e., one entry per row, plus one).

$IA(i)$ contains the index in A of the first nonzero element of row i . Row i of the original matrix extends from $A(IA(i))$ to $A(IA(i+1) - 1)$. The third array, JA , contains the column index of each element of A , so it also is of length NNZ .

For example, the matrix

```
[ 1 2 0 0 ]
[ 0 3 9 0 ]
[ 0 1 4 0 ]
```

is a three-by-four matrix with six nonzero elements, so

```
A = [ 1 2 3 9 1 4 ]
IA = [ 1 3 5 7 ]
JA = [ 1 2 2 3 2 3 ]
```

2. Make yourselves familiar with the YSMF, with generating this specific format (from the “obvious” matrix data structure), and with using it (i.e., accessing non-zero matrix elements). Then, think about the manipulation of the YSMF data structure, in particular what can be readily parallelized, what is less or barely suited for parallelization, and *what* needs to be locked *when*, in case of concurrent, modifying accesses (stores) to the data structure. Then, actually go about parallelizing the problem.

3. Start testing and debugging the program on a single-processor/core computer, then move on to a multi-core or multi-processor system. The matrices A and B should have integer elements and should be randomly initialized, i.e., have random non-zero elements at random positions in the matrices. You should, however, be able to control the degree of density, i.e., the fraction of non-zero elements in the matrices, e.g., 3%. Use very big matrices and YSMF data structures, respectively, up to the limit of storage you can allocate on the test machine; then, parallelization should pay off. The initialization of the matrices and the creation of the YSMF data structures for A and B need *not* be parallelized and *not* be measured.

4. Experiment with a varying number of threads p , e.g., $p = 1, 2, 4, 6, 8, 12, 16$. Determine and compare the *speedup* of the program with varying p . The performance metric is *program run time of the core sparse matrix addition*, i.e., without considering and measuring matrix

initialization and YSMF generation. The reference point, i.e., $speedup = 1$, is the program run time of the core sparse matrix addition with $p = 1$.

5. Work as two- or three-person teams. (Organize yourselves into the teams.) Please send me your solution until 12th May 2016, and do prepare a brief presentation of your approach, solution, and results. There should be one or two presentations on the 12th May; each team should present at least once in the course of the semester.

6. Please use your own (or some other available) multi-core machine; it should have at least two cores.